

# Convolutional Neural Networks (CNN)

## Part1 – Linear Classifiers

# Outline

Part I: Introduction to Visual Recognition



Part II: Image Classification - Challenges

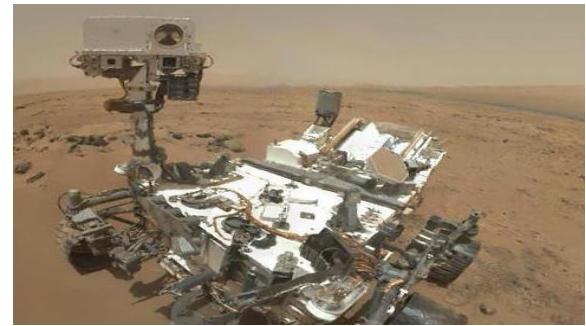


Part III: Image Classification - Example



Part IV: Linear Classifier

# Visual Recognition

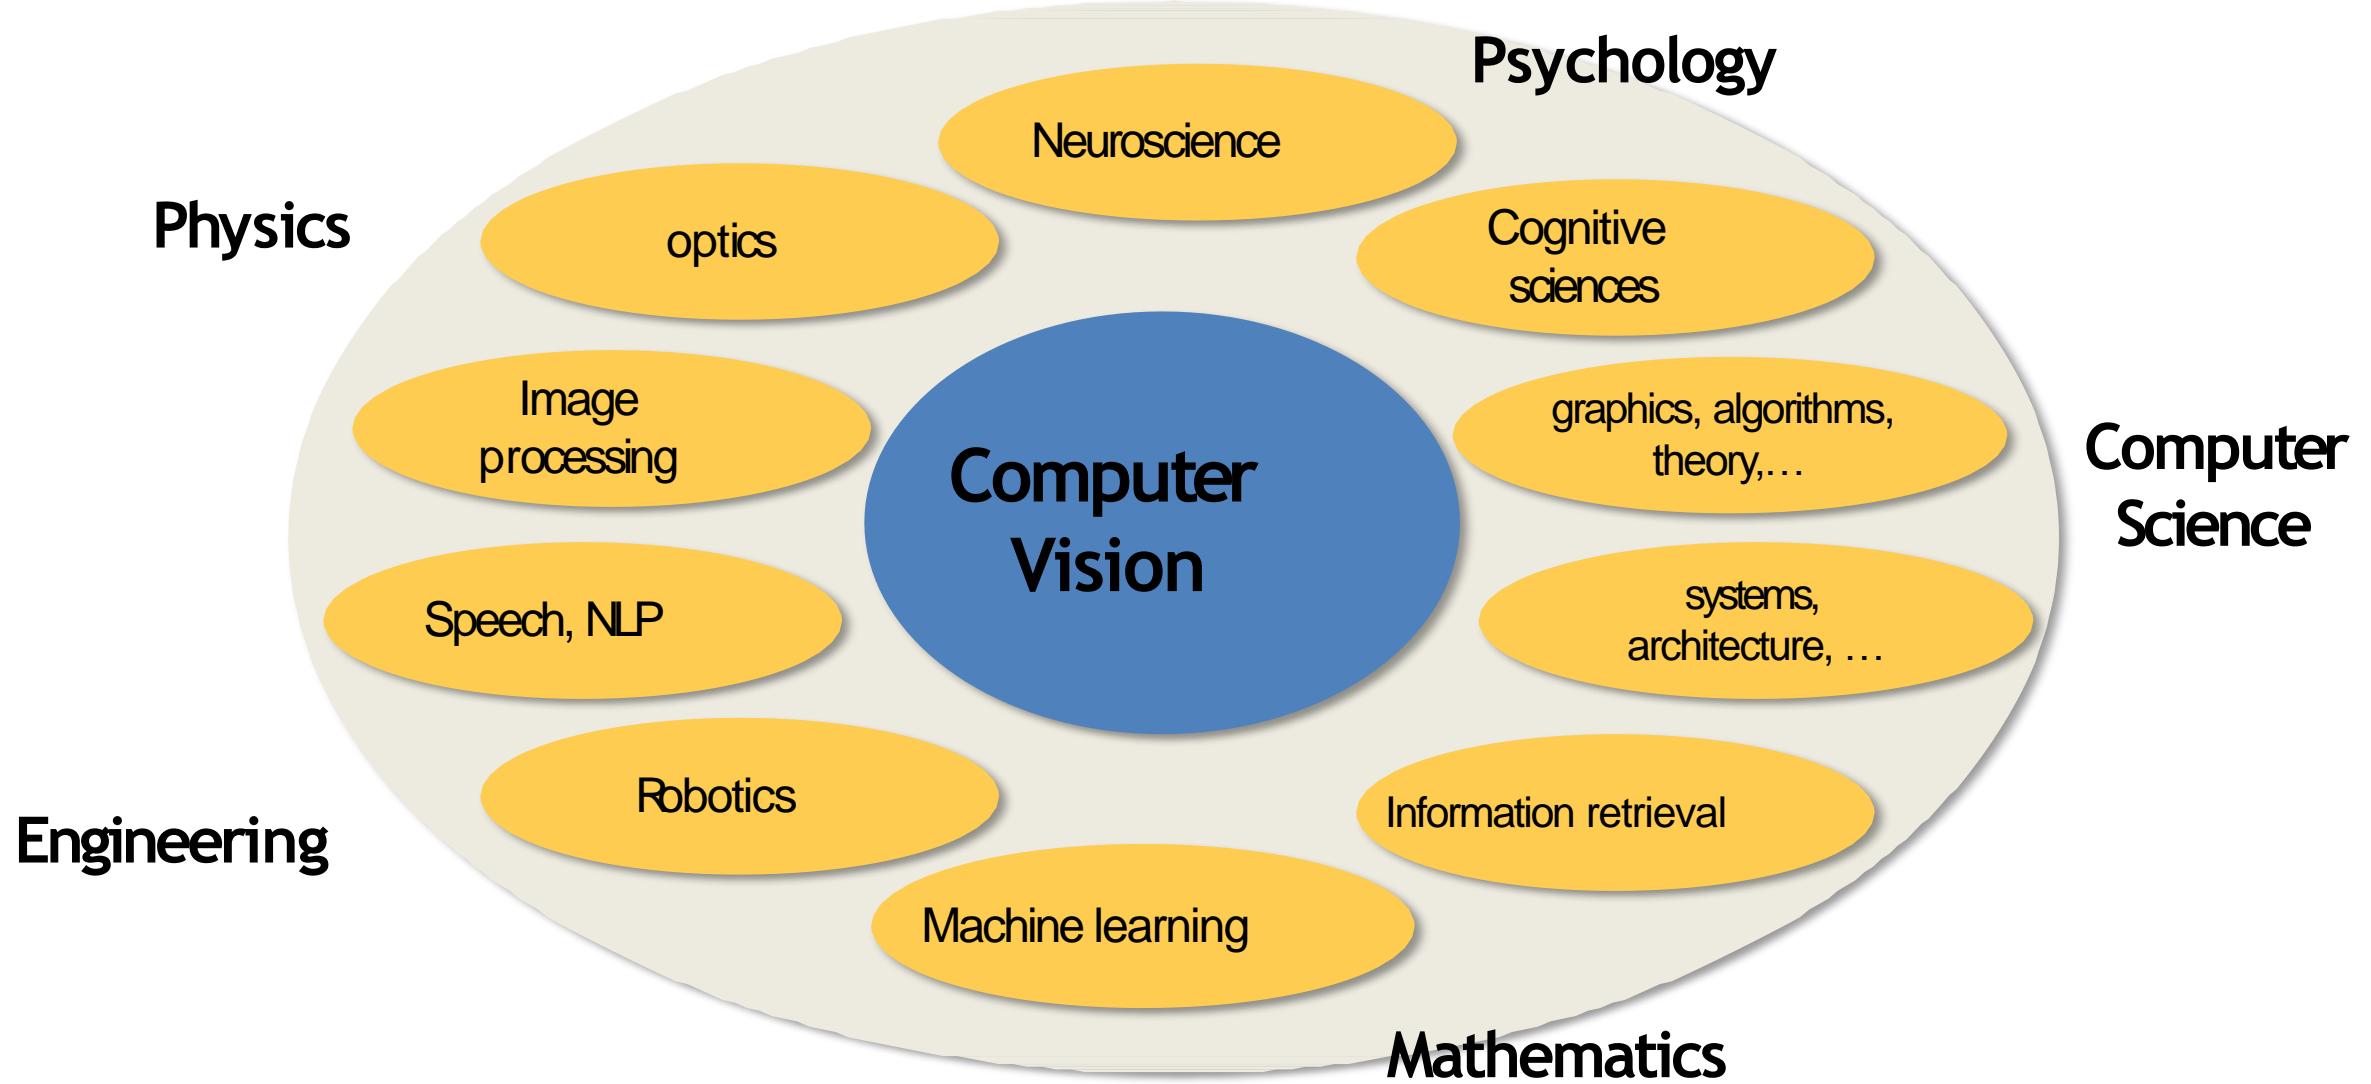


Top row, left to right:  
[Image by Roger H Goun](#) is licensed under [CCBY2.0](#). [Image](#) is [CC01.0](#) public domain  
[Image](#) is [CC0 1.0](#) public domain  
[Image](#) is [CC0 1.0](#) public domain

Middle row, left to right:  
[Image by BGPHConference](#) is licensed under [CCBY2.0](#); changes made. [Image](#) is [CC01.0](#) public domain  
[Image by NASA](#) is licensed under [CCBY2.0](#). [Image](#) is [CC01.0](#) public domain

Bottom row, left to right: [Image](#) is [CC01.0](#), public domain  
[Image by Derek Keats](#) is licensed under [CCBY2.0](#); changes made. [Image](#) is public domain  
[Image](#) is licensed under [CCBY2.0](#); changes made

# Computer Vision



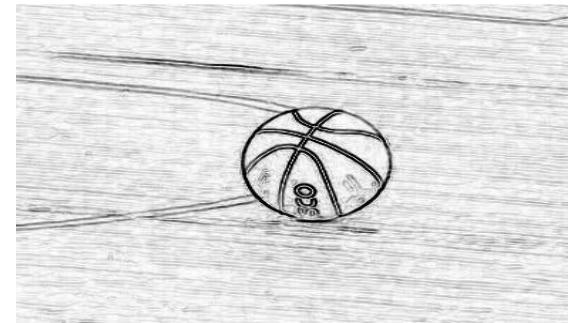
# Stages of Visual Representation, David Marr, 1970s

Input image

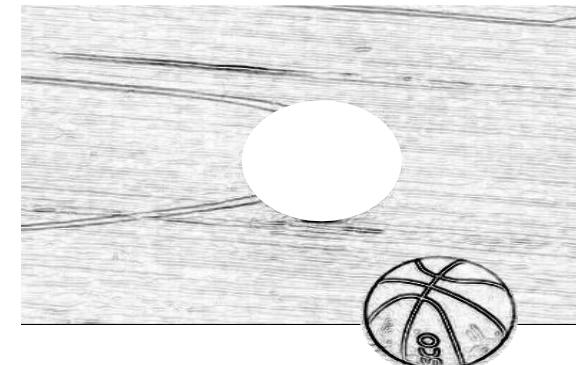


[This image](#) is CC01.0 public domain

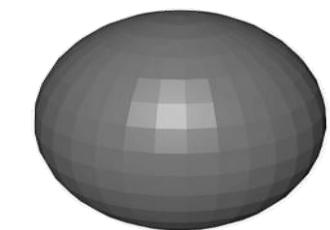
Edge image



2 ½-D SKETCH



3-D model



[This image](#) is CC01.0 public domain

Input  
Image

Primal  
Sketch

2 ½-D  
Sketch

3-D Model  
Representation

Perceived  
intensities

Zero crossings,  
blobs, edges, bars,  
ends, virtual lines,  
groups, curves  
boundaries

Local surface  
orientation and  
discontinuities in  
depth and in surface  
orientation

3-D models  
hierarchically  
organized in terms of  
surface and  
volumetric primitives

Stages of Visual Representation, David Marr, 1970s

# David Lowe, 1987



[Image](#) is CC01.0 public domain



David Lowe, 1987

# Normalized Cut (Shi & Malik, 1997)

Image is CC BY 3.0



Image is public domain



Image is CC BY 2.0,  
changes made



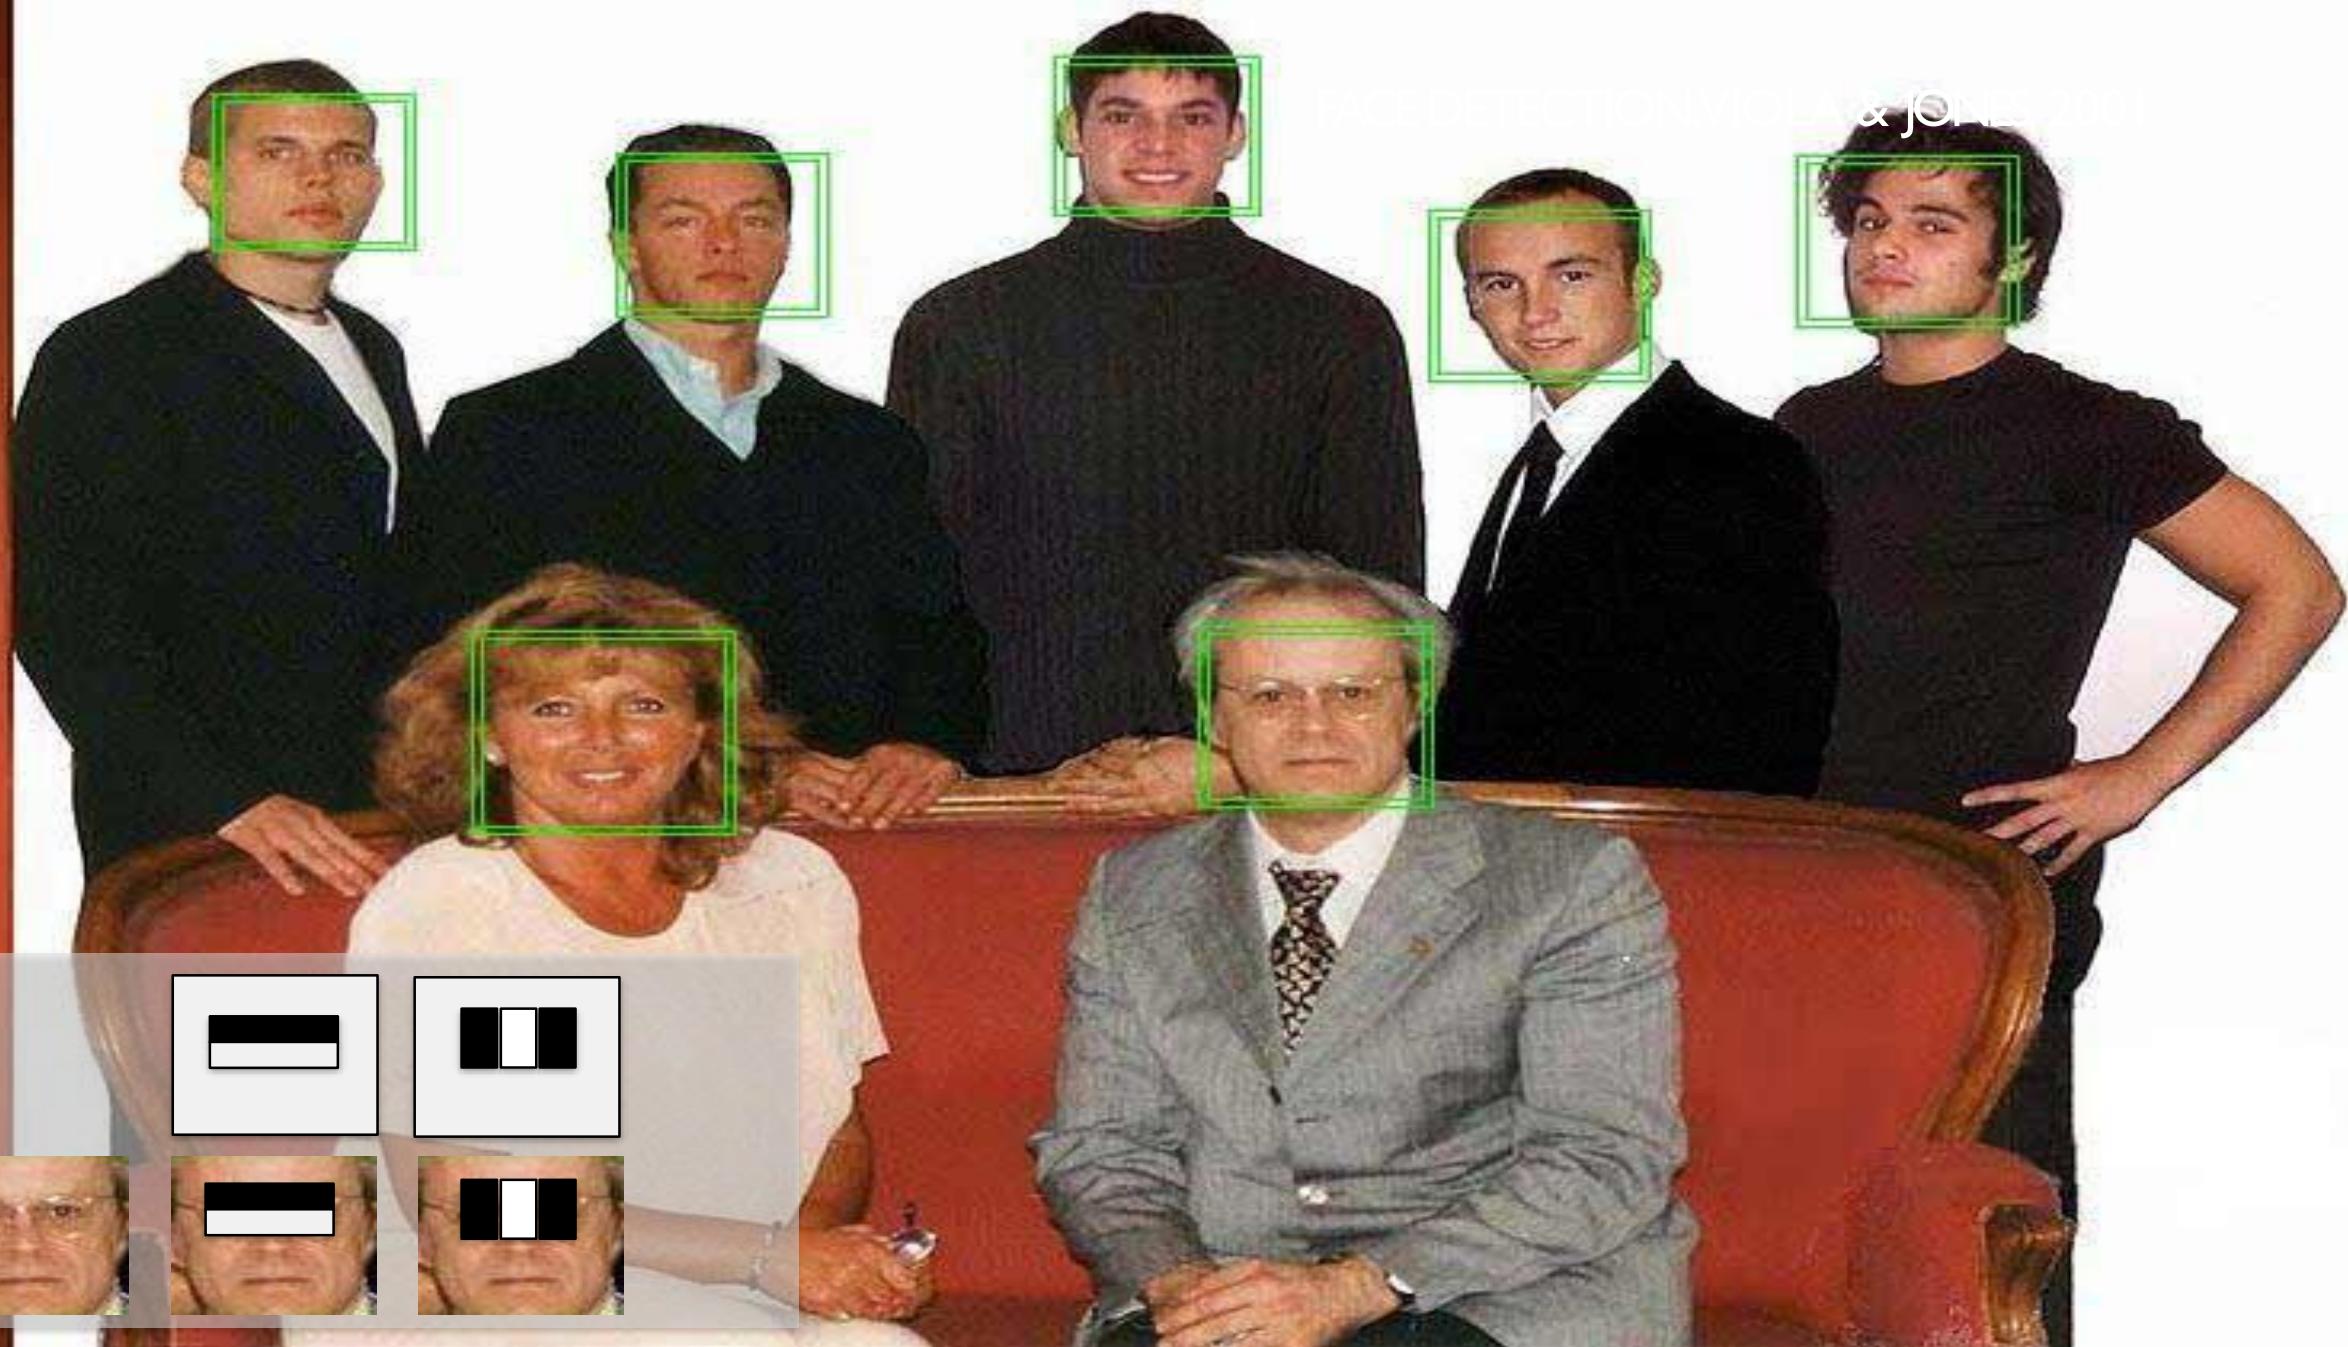


image is public domain

# “SIFT” & Object Recognition, David Lowe, 1999



[Image](#) is public domain



[Image](#) is public domain

SIFT - Scale Invariant Feature Transformation

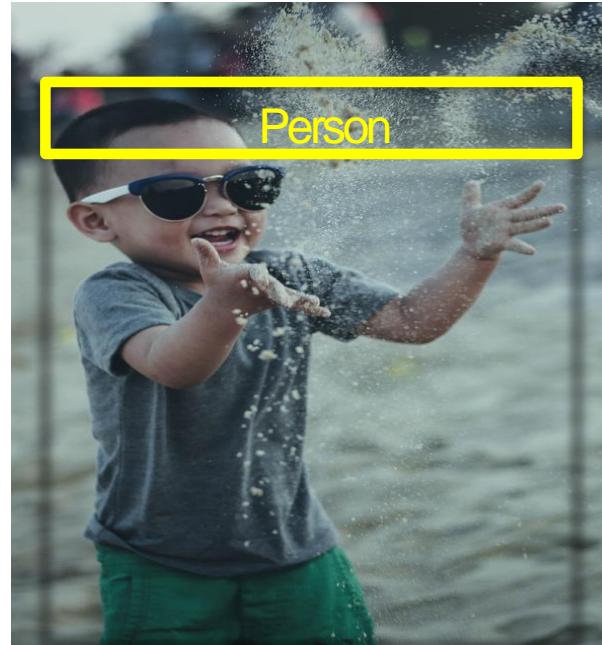
# PASCAL Visual Object Challenge (20 Object Categories)

[Everingham Et Al. 2006-2012]

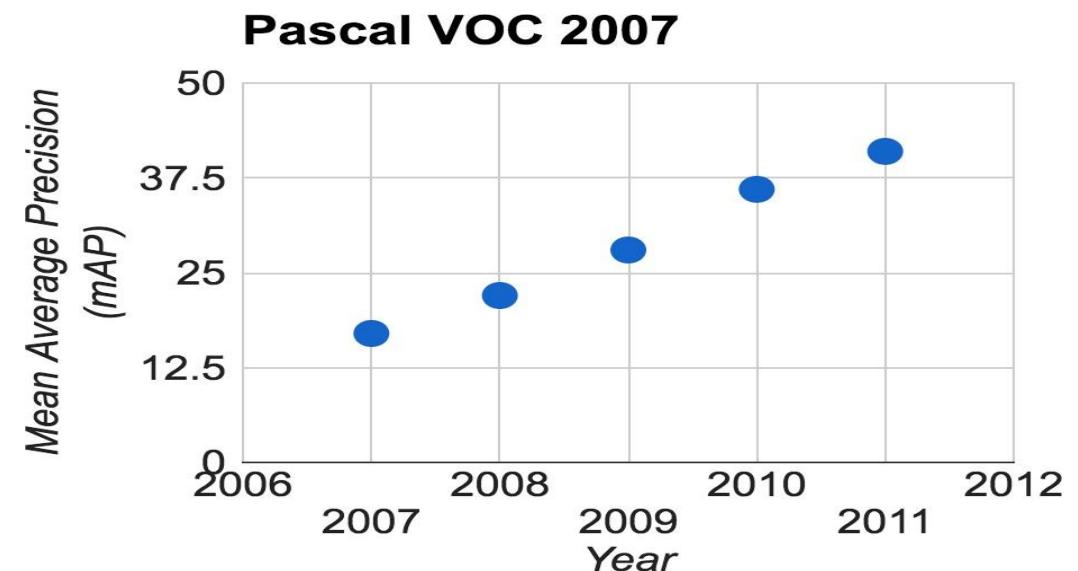
[Image](#) is CC01.0 public domain



[Image](#) is CC01.0 public domain



[Image](#) is CC01.0 public domain





[www.image-net.org](http://www.image-net.org)

## 22K CATEGORIES AND 15M IMAGES

- Animals
  - Bird
  - Fish
  - Mammal
  - Invertebrate
- Plants
  - Tree
  - Flower
- Food
- Materials
- Structures
- Artifact
  - Tools
  - Appliances
  - Structures
- Person
- Scenes
  - Indoor
  - Geological Formations
- Sport Activities

# IMAGENET Large Scale Visual Recognition Challenge

THE IMAGE CLASSIFICATION CHALLENGE  
1,000 OBJECT CLASSES  
120 MILLION IMAGES



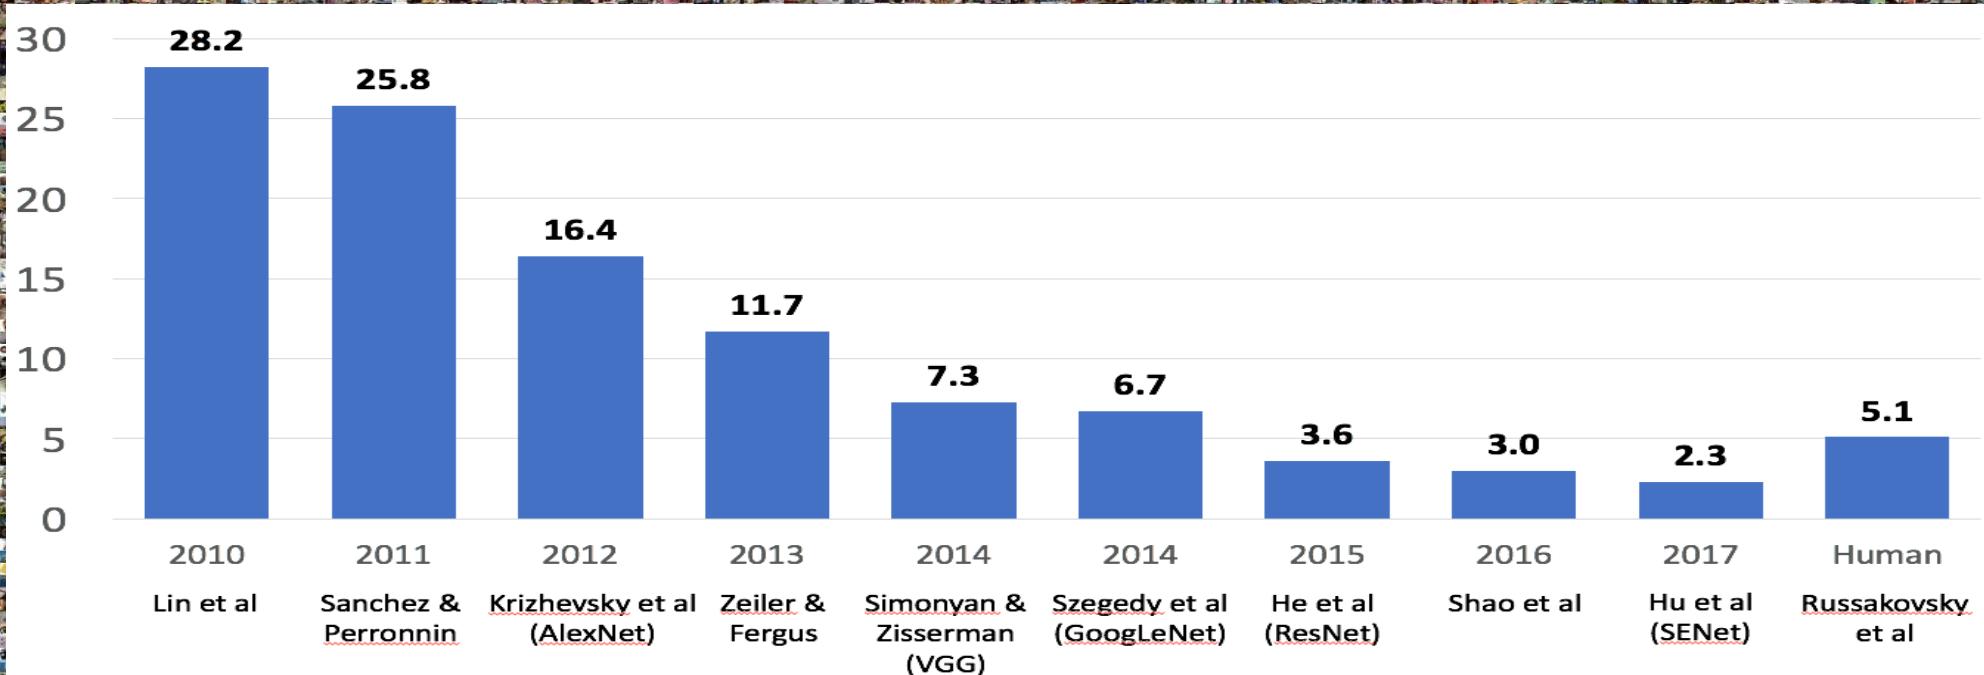
Output:  
Scale

Steel drum  
Drumstick  
Mud turtle



Russakovsky et al. IJCV 2015

The Image Classification Challenge:  
1,000 object classes  
1,431,167 images



Russakovsky et al. IJCV 2015

There Are Many Visual Recognition Problems That Are Related To Image Classification, Such As -

- Object Detection
- Object Recognition
- Image Captioning



This image is licensed under [CCBY-NC-SA 2.0](#); changes made

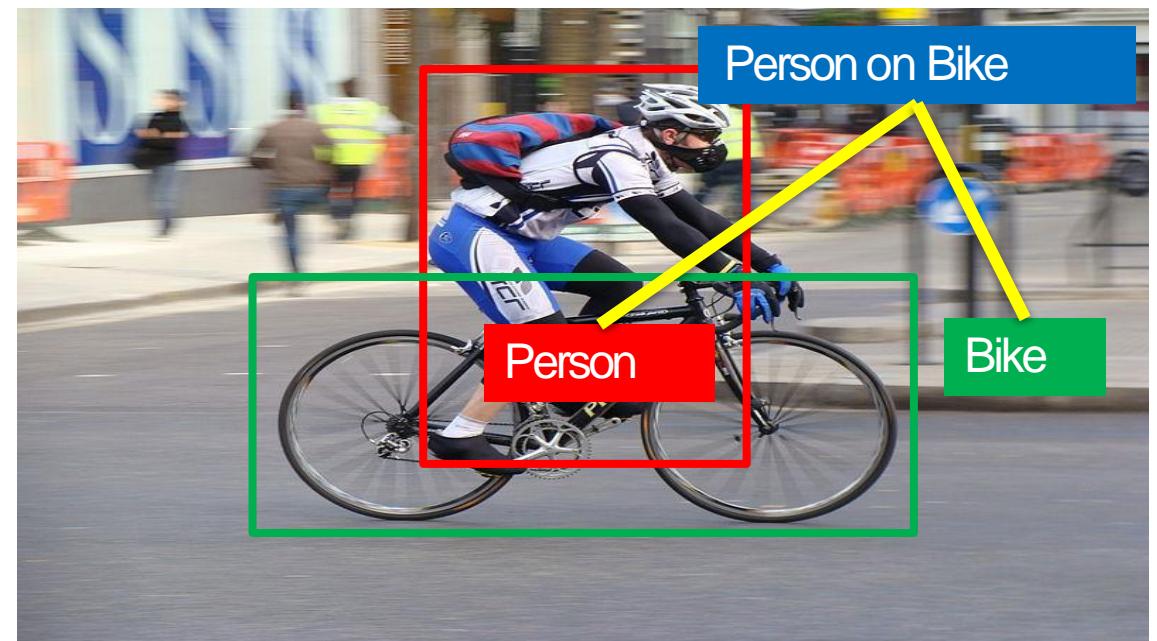


Person

Hammer

This image is licensed under [CCBY-SA 2.0](#); changes made

- Object detection
- Action classification
- Image captioning
- ...

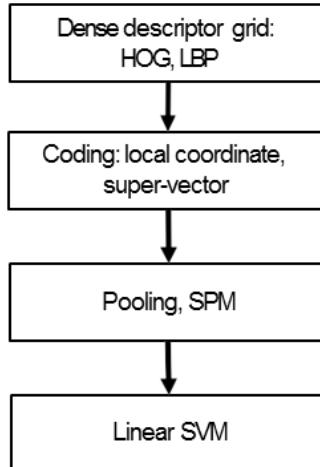


This image is licensed under [CCBY-SA 3.0](#); changes made

# IMAGENET Large Scale Visual Recognition Challenge



NEC-UIUC

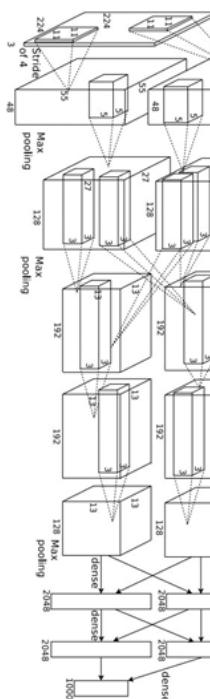


[In CVPR 2011]

Lion image by Swissfrog is  
licensed under [CC-BY 3.0](#)

Year 2012

SuperVision

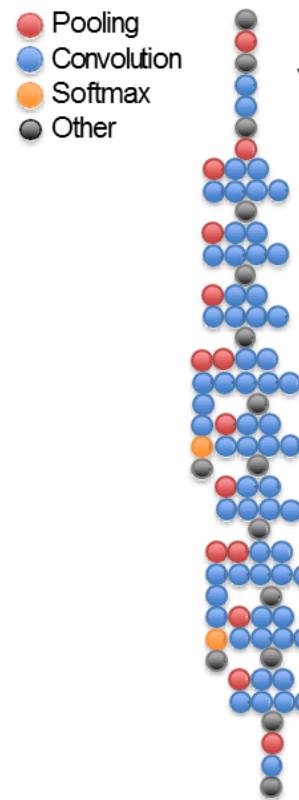


[Krizhevsky NIPS 2012]

Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012.  
Reproduced with permission.

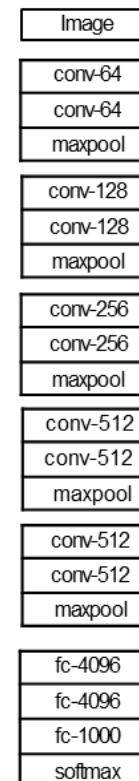
Year 2014

GoogLeNet



[Szegedy arxiv 2014]

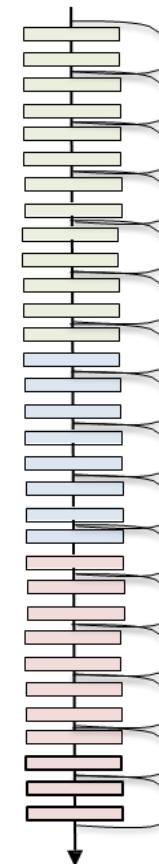
VGG



[Simonyan arxiv 2014]

Year 2015

MSRA



[He ICCV 2015]

CNNs were not invented overnight

# Image Classification

# Image Classification:A core task in Computer Vision



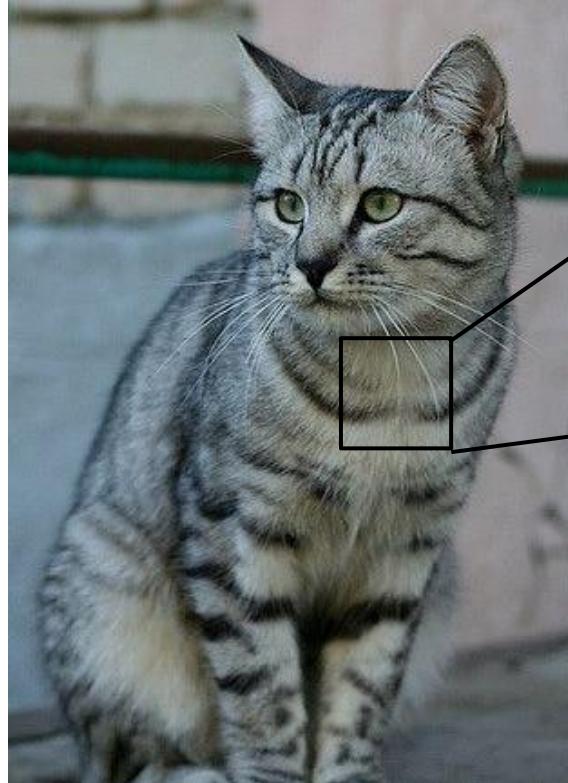
This image by [Nikita](#) is  
licensed under [CC-BY 2.0](#)

(assume given set of discrete labels)  
{dog, cat, truck, plane, ...}



cat

# The Problem: Semantic Gap



[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
[ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
[ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
[ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
[106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
[133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
[128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
[115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
[ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
[ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
[ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
[ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
[ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
[164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
[157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]
[130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
[128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
[123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
[122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]

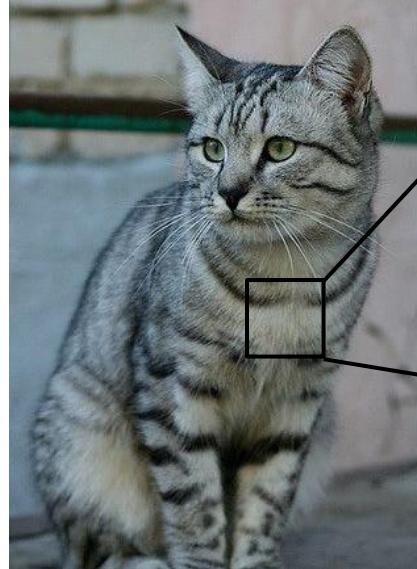
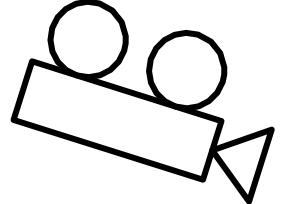
What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3 (3 channels RGB)

This image by [Nikita](#) is  
licensed under [CC-BY 2.0](#)

# Challenges:Viewpoint variation



```
[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
[ 91 98 102 101 104 79 98 103 99 105 123 136 118 105 94 85]
[ 76 85 90 101 128 105 87 96 95 99 115 112 106 103 99 85]
[ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
[106 91 61 63 69 91 88 85 101 107 109 98 75 84 96 95]
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
[133 137 147 103 65 81 88 65 52 54 74 84 102 93 85 82]
[128 137 144 140 109 95 86 70 62 65 63 63 68 73 86 101]
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
[115 114 109 123 153 148 131 118 113 109 100 92 74 65 72 78]
[ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
[ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
[ 62 65 82 89 78 71 89 101 124 126 119 101 107 114 131 119]
[ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
[ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
[164 146 112 86 82 120 124 104 76 48 45 66 88 101 102 109]
[157 170 157 120 93 86 114 132 112 97 69 55 78 82 99 94]
[138 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
[128 112 96 117 153 144 128 115 104 107 102 93 87 81 72 79]
[123 107 96 88 83 112 153 149 122 109 104 75 88 107 112 99]
[122 121 102 88 82 86 94 117 145 148 153 102 58 78 92 107]
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```

All pixels change  
when the camera  
moves!

# Challenges: Background Clutter



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

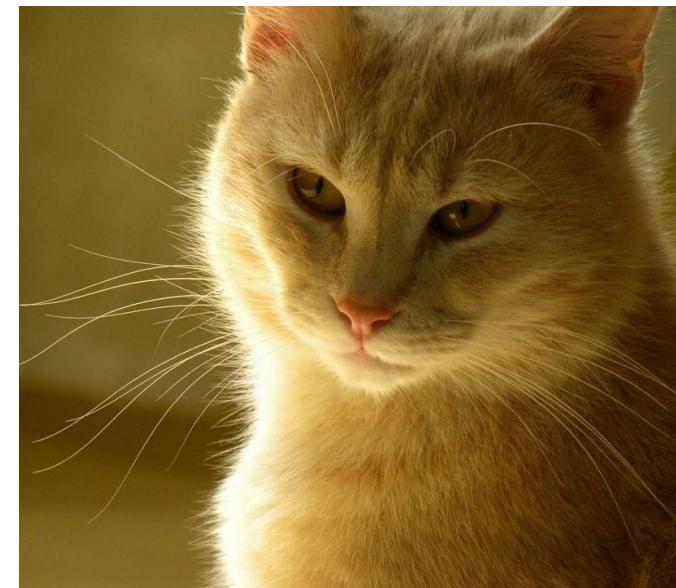
# Challenges: Illumination



[This image](#) is CC0 1.0 public domain



[This image](#) is CC0 1.0 public domain

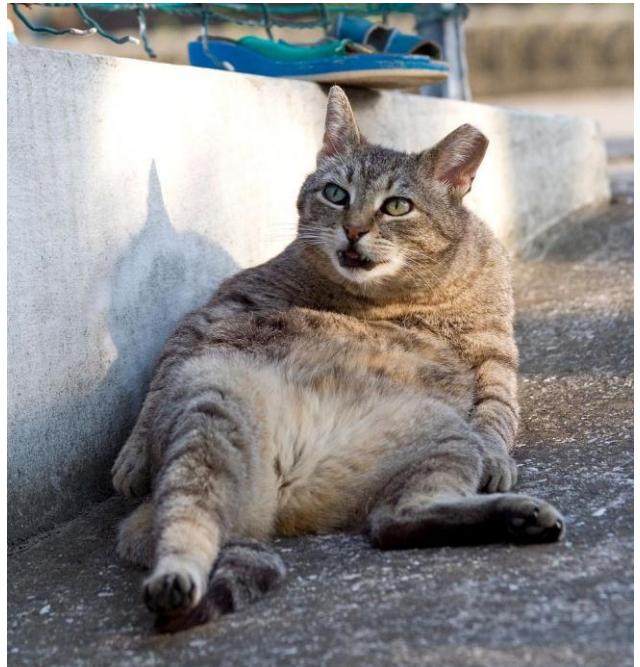


[This image](#) is CC0 1.0 public domain



[This image](#) is CC0 1.0 public domain

# Challenges: Deformation



[This image](#) by [Umberto Salvagnin](#)  
is licensed under [CC-BY 2.0](#)



[This image](#) by [Umberto Salvagnin](#)  
is licensed under [CC-BY 2.0](#)



[This image](#) by [sare bear](#) is  
licensed under [CC-BY 2.0](#)



[This image](#) by [Tom Thai](#) is  
licensed under [CC-BY 2.0](#)

# Challenges: Occlusion



[This image](#) is [CC0 1.0](#) public domain



[This image](#) is [CC0 1.0](#) public domain



[This image](#) by [jonsson](#) is licensed  
under [CC-BY 2.0](#)

# Challenges: Intraclass variation



[This image](#) is CC0 1.0 public domain

# An image classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

Unlike e.g. sorting a list of numbers,

**no obvious way** to hard-code the algorithm for  
recognizing a cat, or other classes.

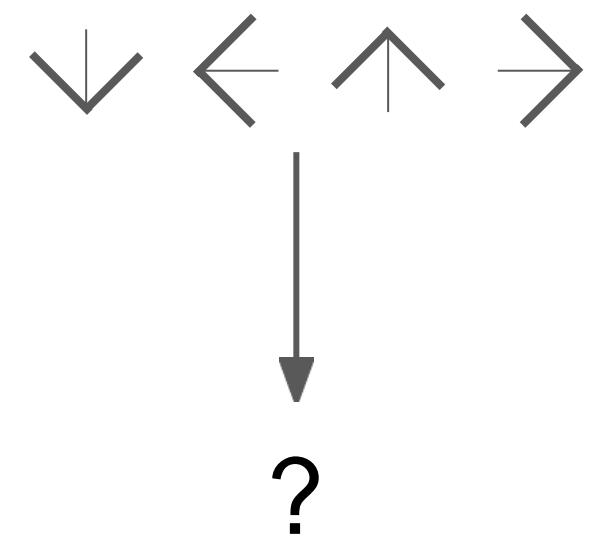
# Attempts have been made



Find edges



Find corners



# Machine Learning: Data-Driven Approach

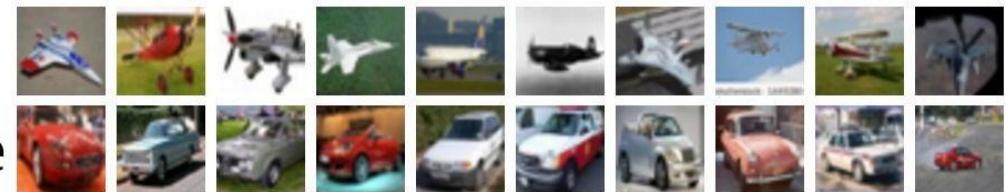
1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images

Example training set

```
def train(images, labels):  
    # Machine learning!  
    return model
```

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```

**airplane**



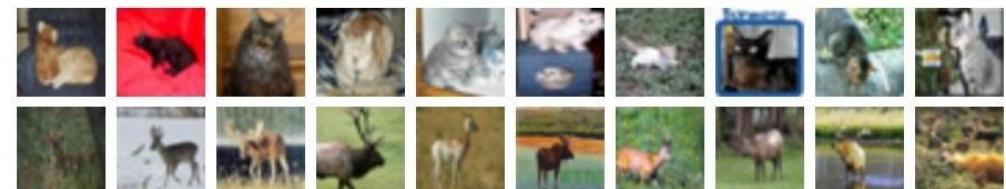
**automobile**



**bird**



**cat**



**deer**



# First classifier: Nearest Neighbor

```
def train(images, labels):  
    # Machine learning!  
    return model
```



Memorize all  
data and labels

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```



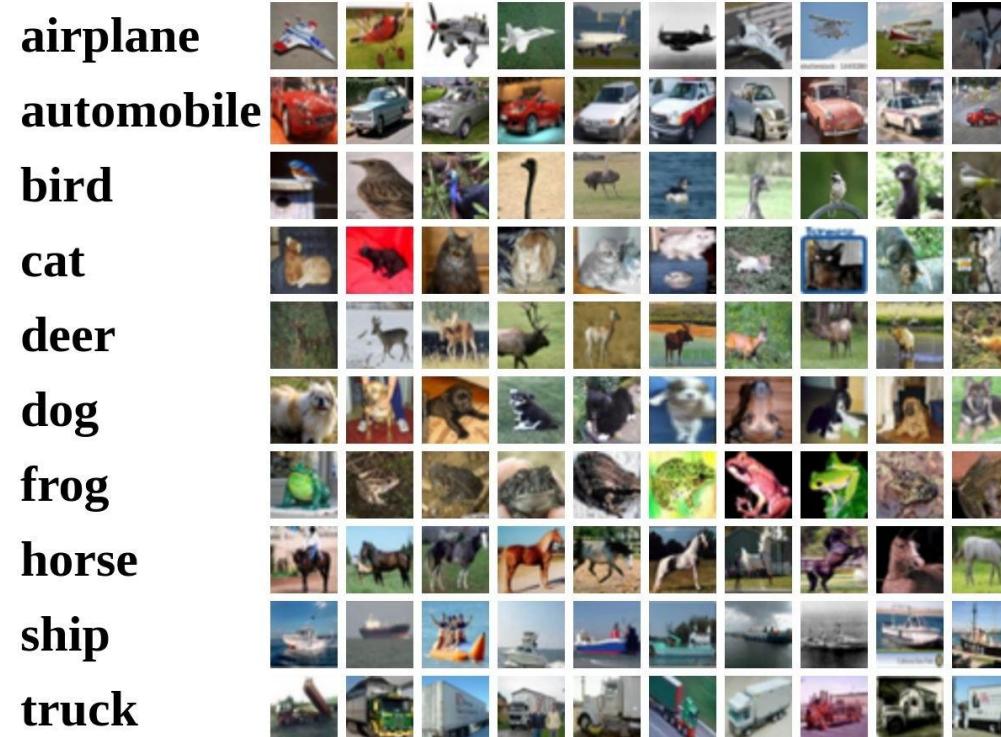
Predict the label  
of the most similar  
training image

# Example Dataset: CIFAR10

**10 classes**

**50,000 training images**

**10,000 testing images**

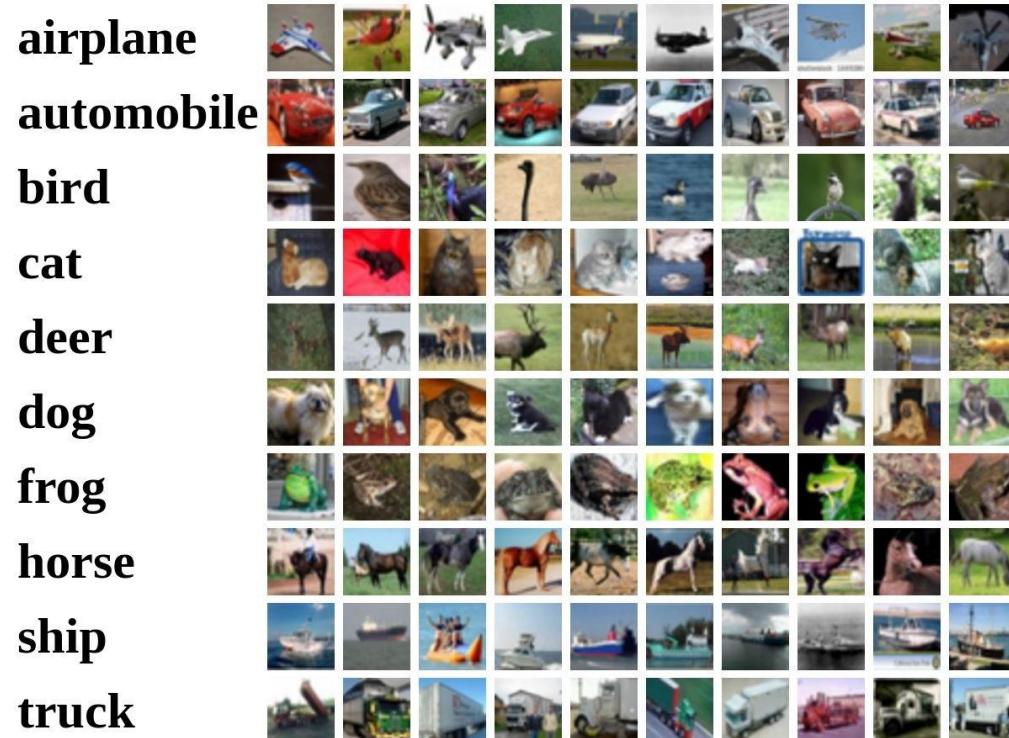


# Example Dataset: CIFAR10

10 classes

50,000 training images

10,000 testing images



Test images and nearest neighbors



# Distance Metric to compare images

**L1 distance:**

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image

56	32	10	18
90	23	128	133
24	26	178	200
2	0	255	220

training image

10	20	24	17
8	10	89	100
12	16	178	170
4	32	233	112

pixel-wise absolute value differences

46	12	14	1
82	13	39	33
12	10	0	30
2	32	22	108

add → 456

# Nearest Neighbor classifier

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

    return Ypred
```

## Nearest Neighbor classifier

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Memorize training data

# Nearest Neighbor Classifier

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

For each test image:  
Find closest train image  
Predict label of nearest image

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

## Nearest Neighbor classifier

**Q: With N examples, how fast are training and prediction?**

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

## Nearest Neighbor classifier

**Q:** With  $N$  examples, how fast are training and prediction?

**A:** Train  $O(1)$ , predict  $O(N)$

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

## Nearest Neighbor classifier

**Q:** With  $N$  examples, how fast are training and prediction?

**A:** Train  $O(1)$ , predict  $O(N)$

This is bad: we want classifiers that are **fast** at prediction; **slow** for training is ok

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

## Nearest Neighbor Classifier

Many methods exist for fast / approximate nearest neighbor (beyond the scope of 231N!)

A good implementation:

<https://github.com/facebookresearch/faiss>

Johnson et al, “Billion-scale similarity search with GPUs”, arXiv 2017

# What does this look like?



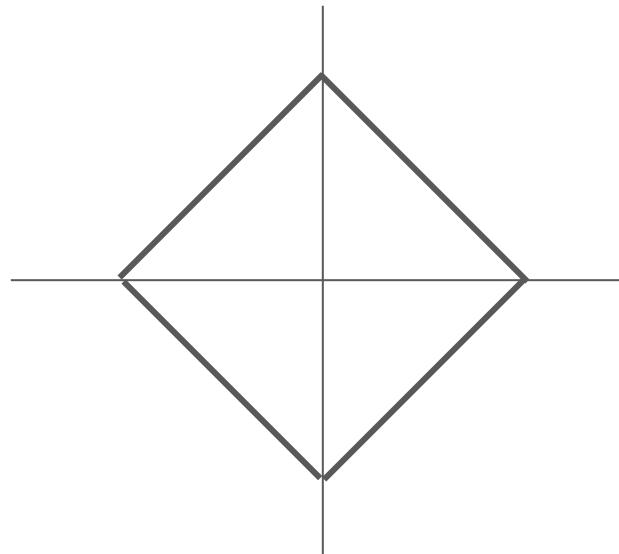
# What does this look like?



# K-Nearest Neighbors: Distance Metric

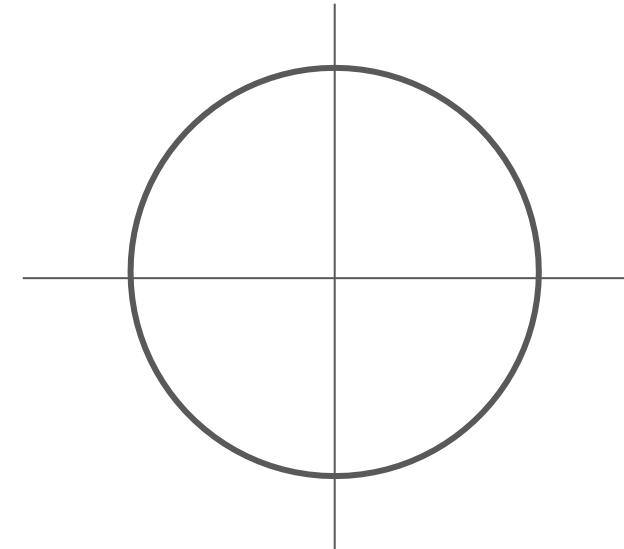
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

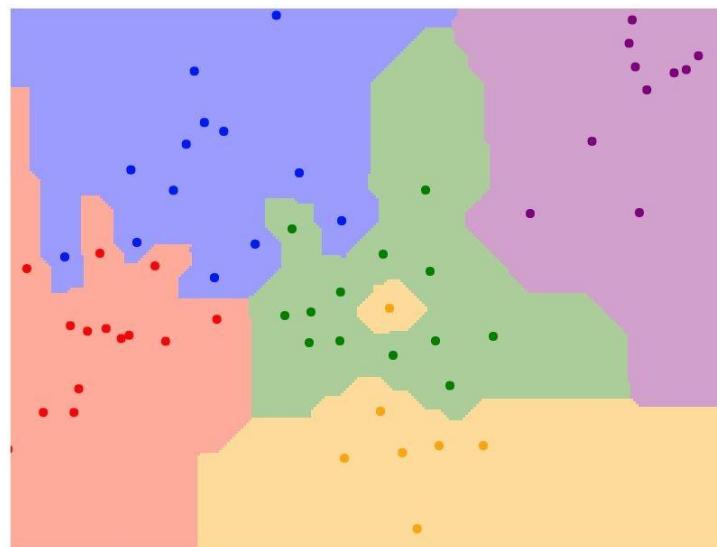
$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



# K-Nearest Neighbors: Distance Metric

L1 (Manhattan) distance

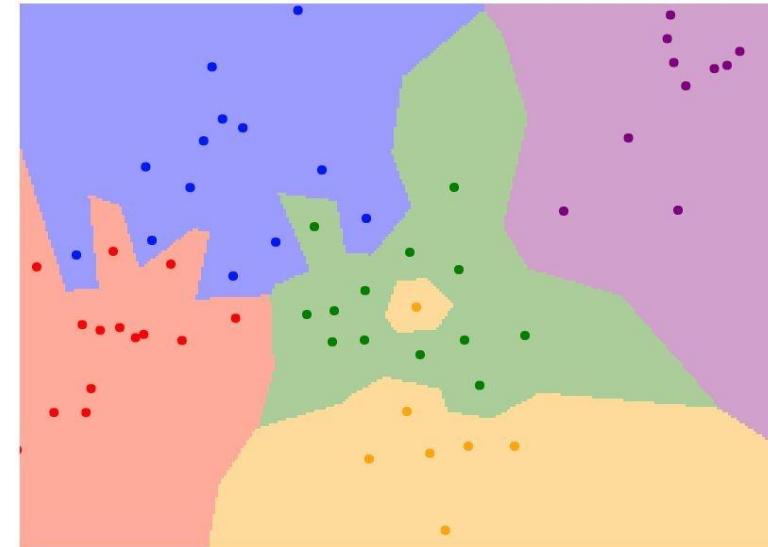
$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



$K = 1$

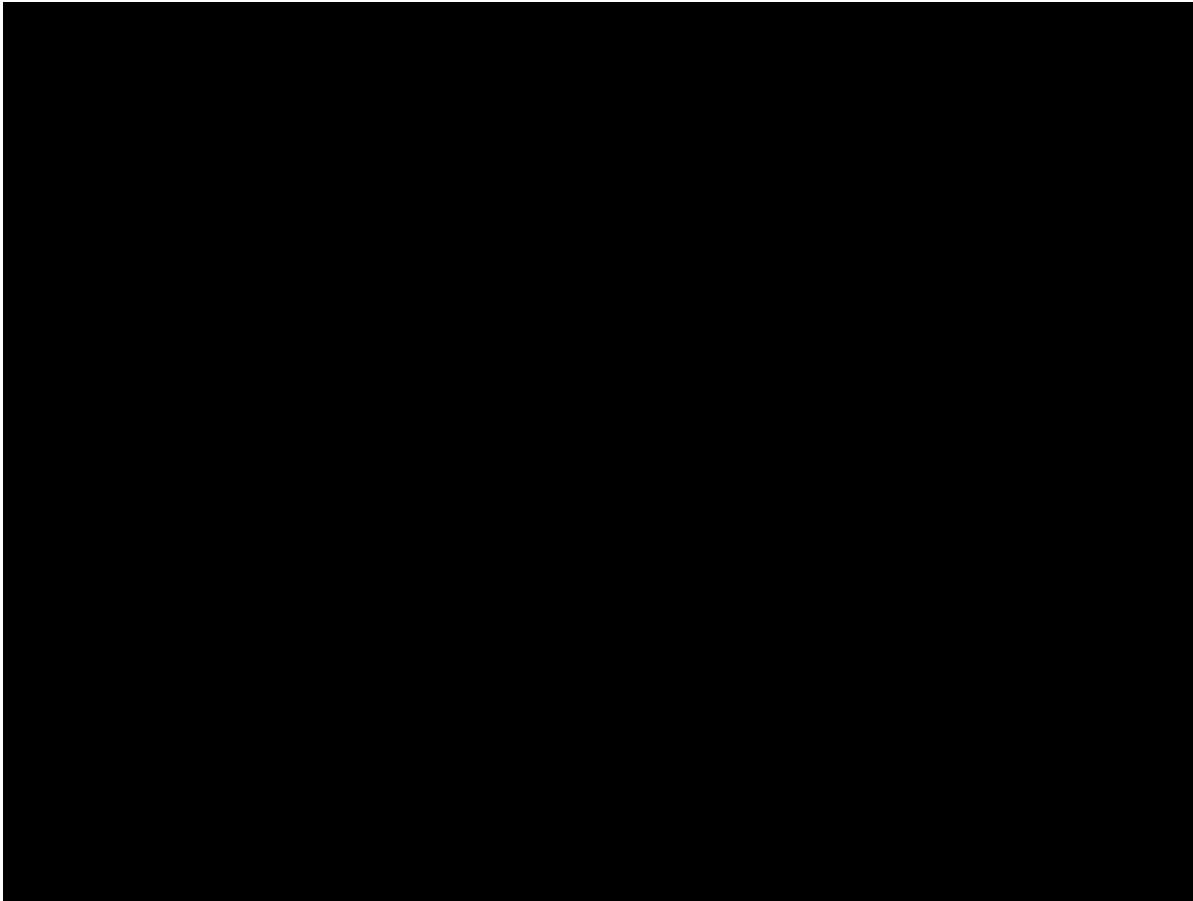
L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



$K = 1$

# K-Nearest Neighbors: Demo



<http://vision.stanford.edu/teaching/cs231n-demos/knn/>

# Hyperparameters

- What is the best value of  $k$  to use? What is the best distance to use?
- These are **hyperparameters**: choices about the algorithm that we set rather than learn

# Hyperparameters

- What is the best value of  $k$  to use? What is the best distance to use?
- These are hyperparameters: choices about the algorithm that we set rather than learn
- Very problem-dependent.
- Must try them all out and see what works best.

# Setting Hyperparameters

**Idea #1:** Choose hyperparameters that work best on the data

Your Dataset

# Setting Hyperparameters

**Idea #1:** Choose hyperparameters  
that work best on the data

**BAD:** K = 1 always works  
perfectly on training data

Your Dataset

# Setting Hyperparameters

**Idea #1:** Choose hyperparameters that work best on the data

**BAD:** K = 1 always works perfectly on training data

Your Dataset

**Idea #2:** Split data into **train** and **test**, choose hyperparameters that work best on test data

train

test

# Setting Hyperparameters

**Idea #1:** Choose hyperparameters that work best on the data

**BAD:** K = 1 always works perfectly on training data

Your Dataset

**Idea #2:** Split data into **train** and **test**, choose hyperparameters that work best on test data

train

**BAD:** No idea how algorithm will perform on new data

test

# Setting Hyperparameters

**Idea #1:** Choose hyperparameters that work best on the data

**BAD:** K = 1 always works perfectly on training data

Your Dataset

**Idea #2:** Split data into **train** and **test**, choose hyperparameters that work best on test data

train

**BAD:** No idea how algorithm will perform on new data

test

**Idea #3:** Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

**Better!**

train

validation

test

# Setting Hyperparameters

Your Dataset

**Idea #4: Cross-Validation:** Split data into **folds**,  
try each fold as validation and average the results

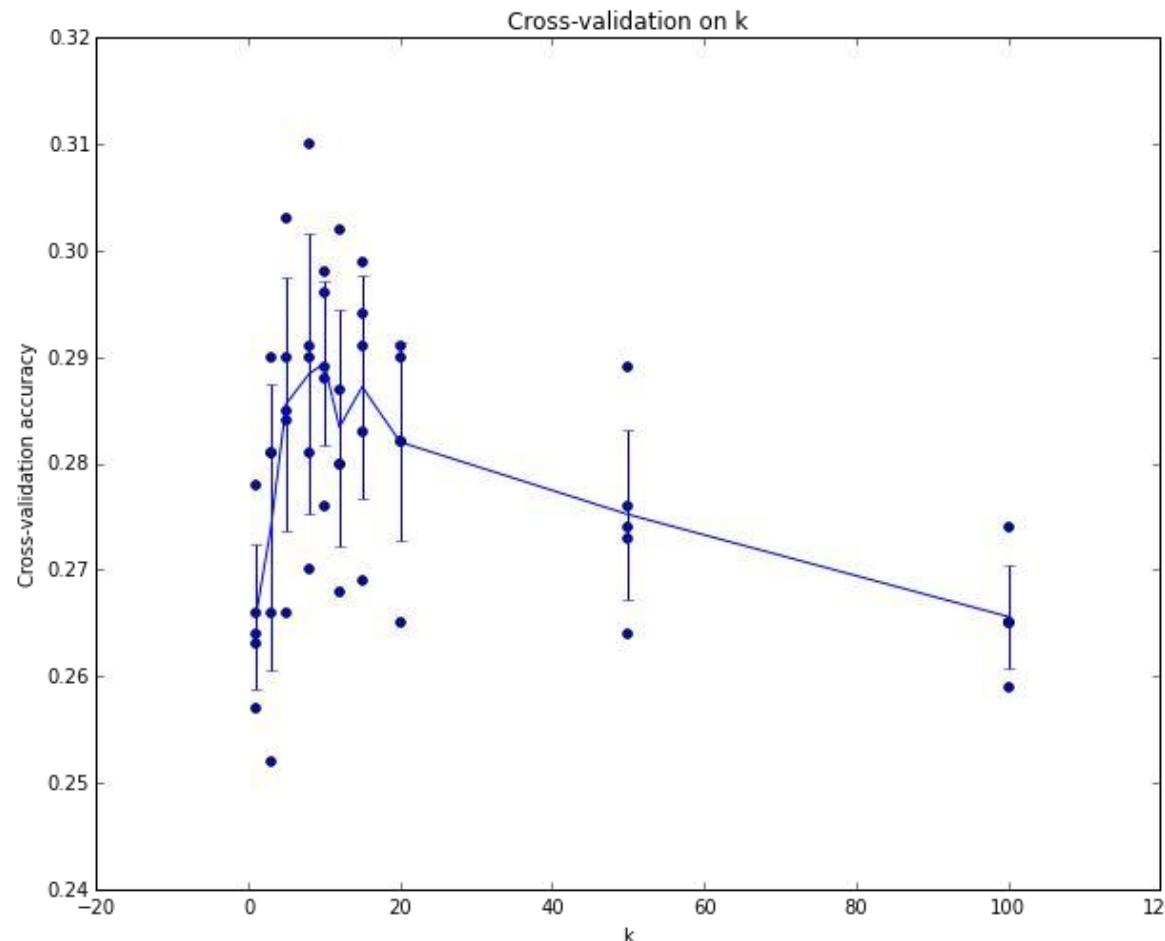
fold 1	fold 2	fold 3	fold 4	fold 5	test
--------	--------	--------	--------	--------	------

fold 1	fold 2	fold 3	fold 4	fold 5	test
--------	--------	--------	--------	--------	------

fold 1	fold 2	fold 3	fold 4	fold 5	test
--------	--------	--------	--------	--------	------

Useful for small datasets, but not used too frequently in deep learning

# Setting Hyperparameters



Example of  
5-fold cross-validation  
for the value of **k**.

Each point: single  
outcome.

The line goes  
through the mean, bars  
indicated standard  
deviation

(Seems that  $k \sim 7$  works best  
for this data)

# k-Nearest Neighbor on images never used.

- Very slow at test time
- Distance metrics on pixels are not informative

Original



Boxed



Shifted



Tinted



(all 3 images have same L2 distance to the one on the left)

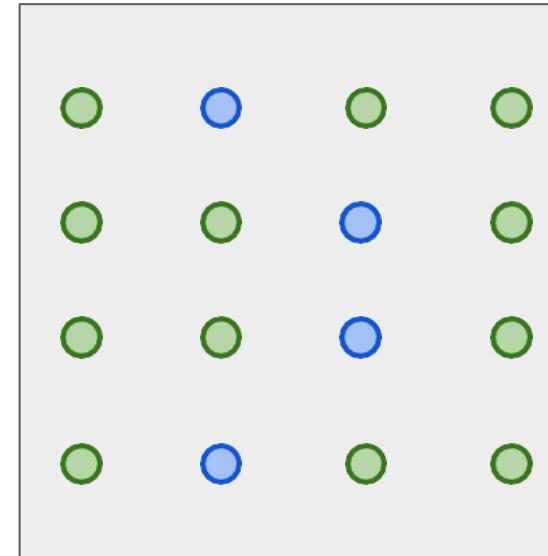
# k-Nearest Neighbor on images never used.

- Curse of dimensionality

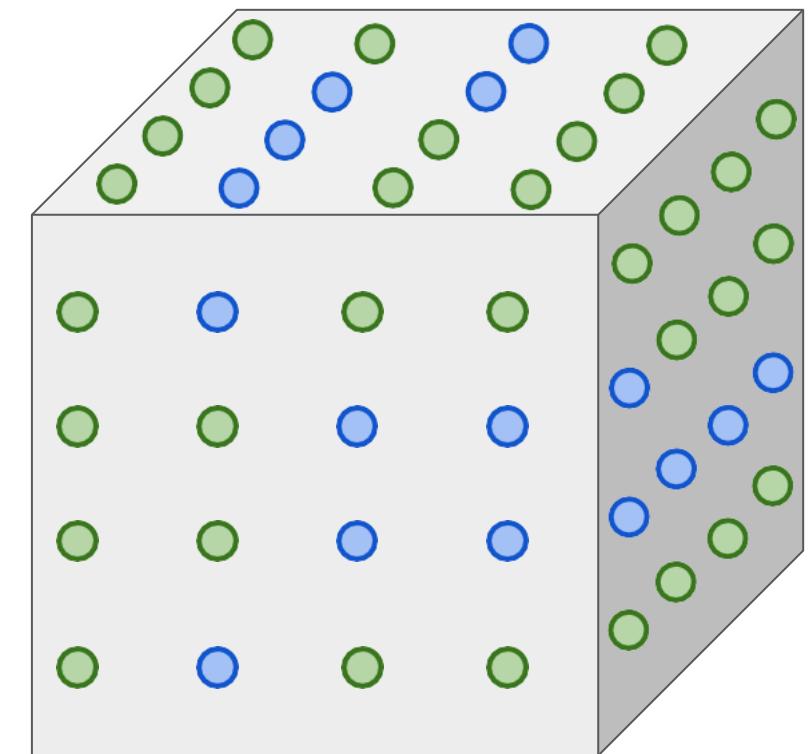
Dimensions = 1  
Points = 4



Dimensions = 2  
Points =  $4^2$



Dimensions = 3  
Points =  $4^3$



# K-Nearest Neighbors: Summary

In **Image classification** we start with a **training set** of images and labels, and must predict labels on the **test set**

The **K-Nearest Neighbors** classifier predicts labels based on nearest training examples

Distance metric and K are **hyperparameters**

Choose hyperparameters using the **validation set**;  
only run on the test set once at the very end!

# Linear Classification

# Neural Network

Linear  
classifiers



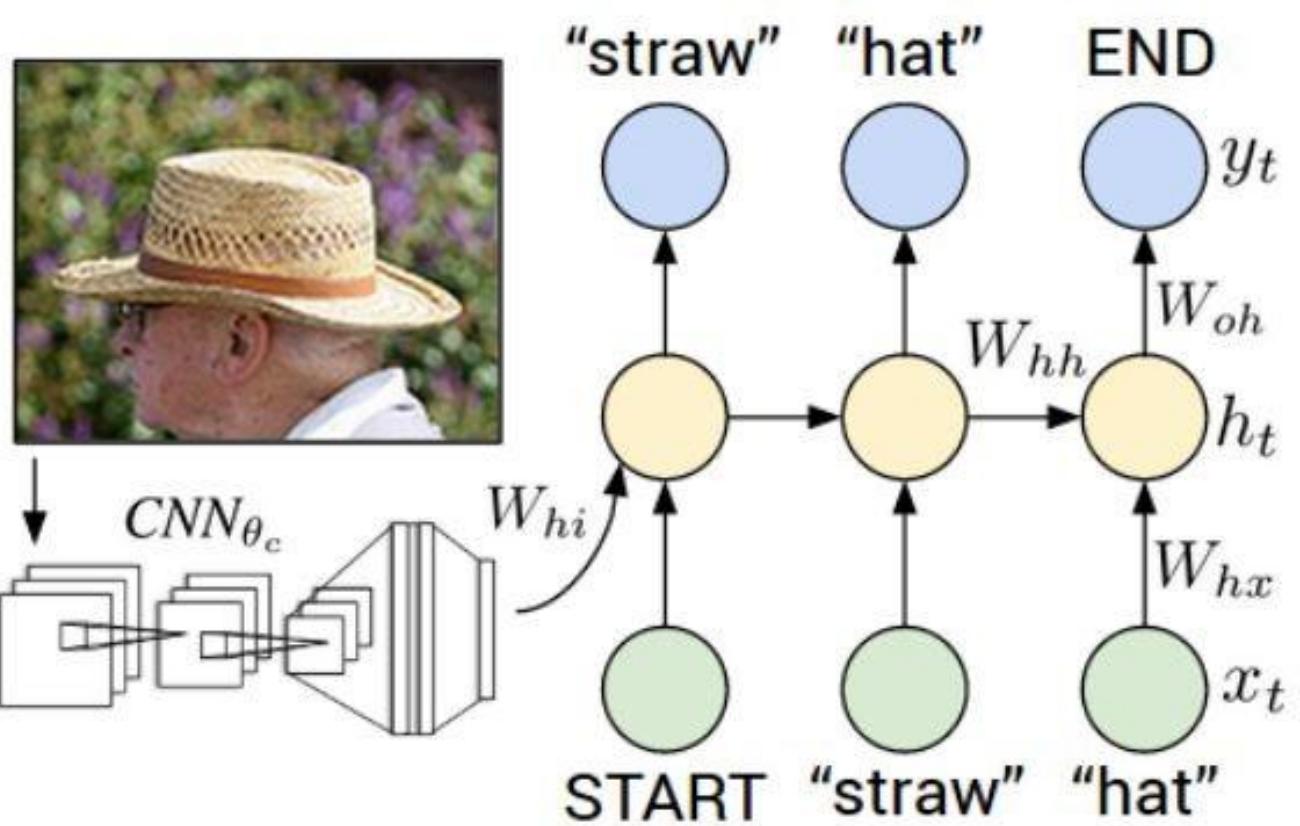
*Two young girls are playing with lego toy. on wakeboard*



*Boy is doing backflip*



*Construction worker in orange safety vest is working on road.*



Karpathy and Fei-Fei, "Deep Visual-Semantic Alignments for Generating Image Descriptions", CVPR 2015  
Figures copyright IEEE, 2015. Reproduced for educational purposes.

# Recall CIFAR10

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



**50,000** training images  
each image is **32x32x3**

**10,000** test images.

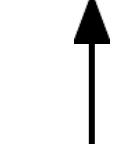
# Parametric Approach



Array of **32x32x3** numbers  
(3072 numbers total)

Image

$$f(\mathbf{x}, \mathbf{W})$$



**W**  
**parameters  
or weights**

10 numbers giving  
class scores

# Parametric Approach: Linear Classifier

$$F(X, W) = WX$$

Image



Array of **32x32x3** numbers  
(3072 numbers total)

$$f(x, W)$$



**W**  
**parameters**  
**or weights**

10 numbers giving  
class scores

# Parametric Approach: Linear Classifier

Image



Array of **32x32x3** numbers  
(3072 numbers total)

$$f(x, W) = \boxed{W} \boxed{x}$$

**10x1      10x3072**

$$f(\textcolor{blue}{x}, \textcolor{red}{W})$$



**W**  
**parameters**  
**or weights**

**3072x1**

10 numbers giving  
class scores

# Parametric Approach: Linear Classifier

Image



Array of **32x32x3** numbers  
(3072 numbers total)

$$f(x, W) = \boxed{W} \boxed{x} + \boxed{b} \quad 10 \times 1$$

**10x1      10x3072**

$f(\boxed{x}, \boxed{W})$

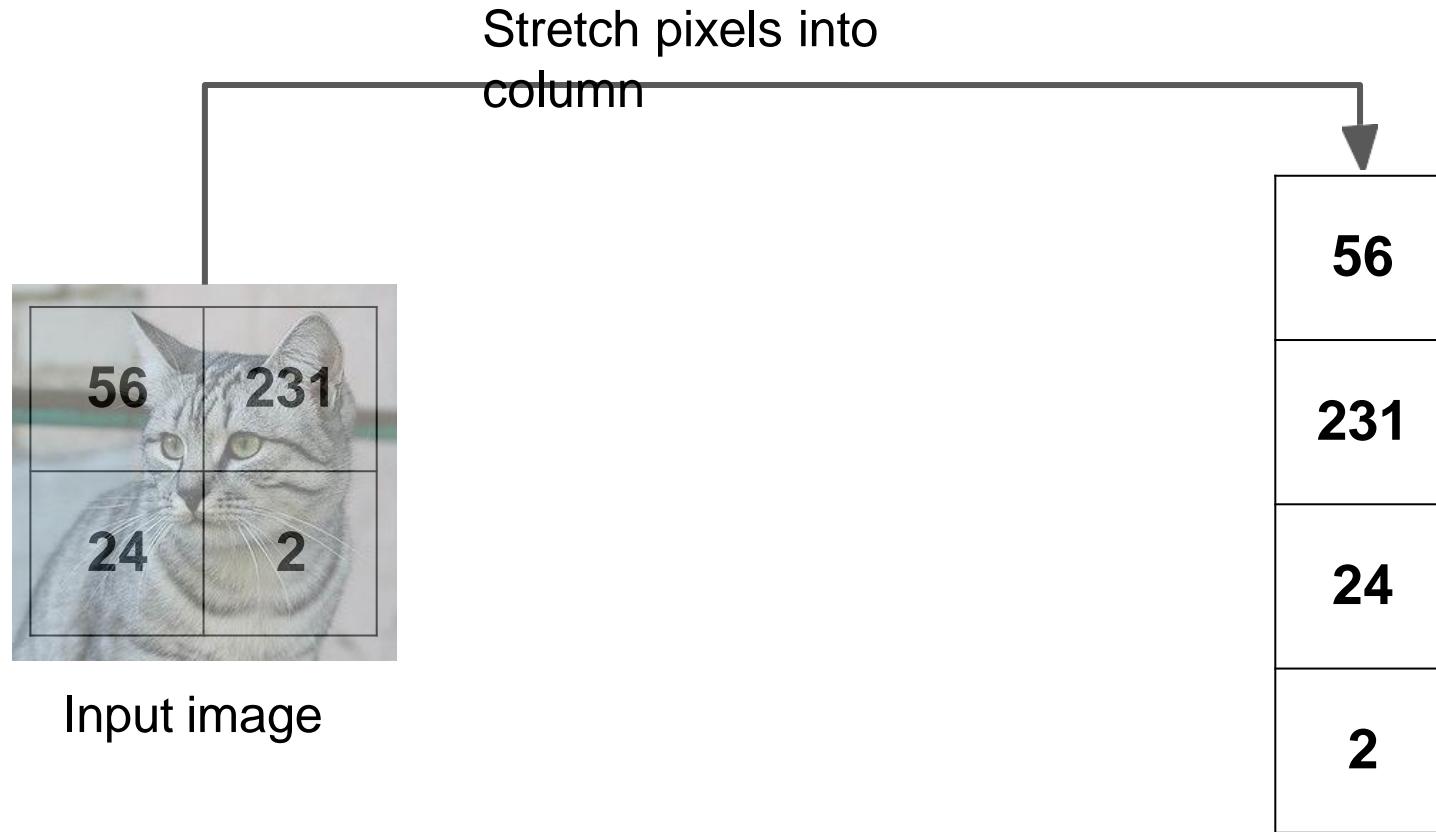
10 numbers giving  
class scores



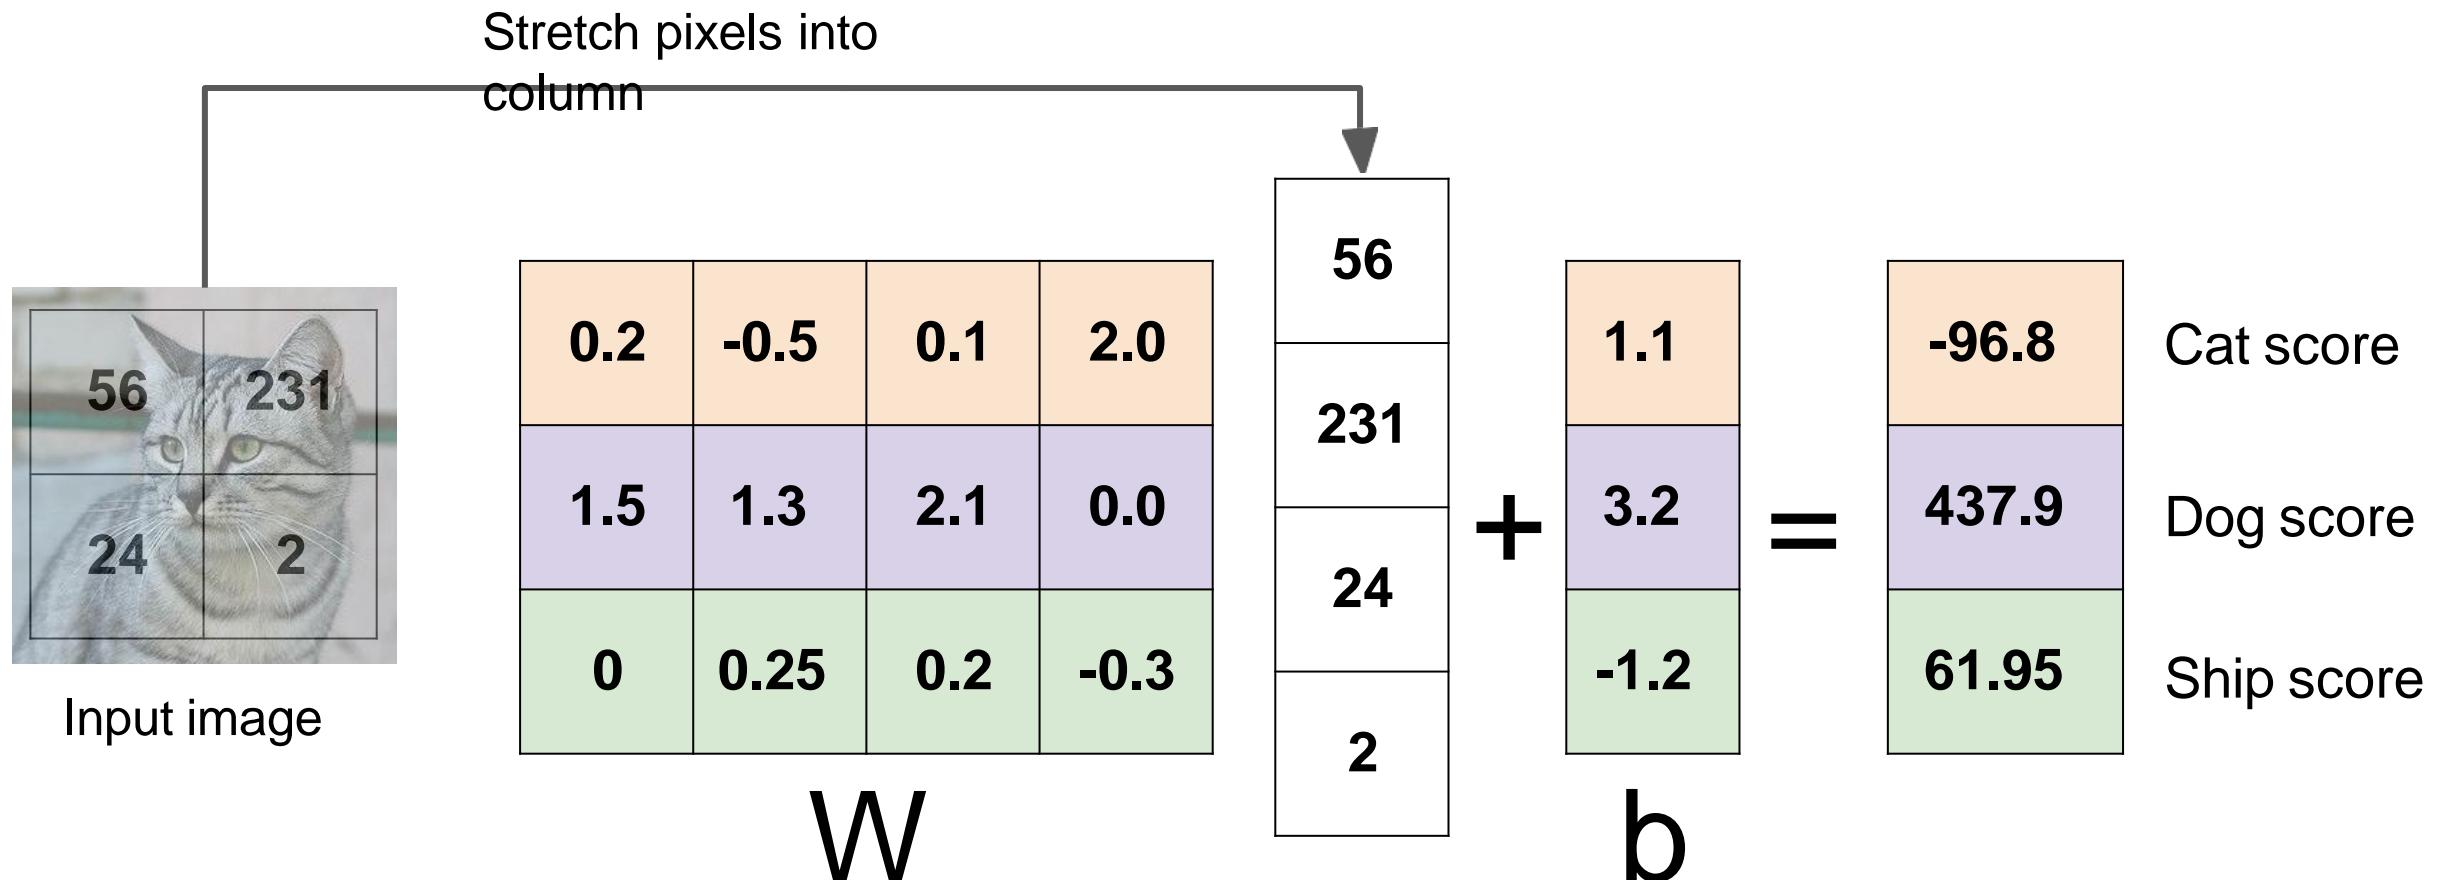
**W**

parameters  
or weights

# Example With An Image With 4 Pixels, And 3 Classes (Cat/Dog/Ship)



# Example With An Image With 4 Pixels, And 3 Classes (Cat/Dog/Ship)



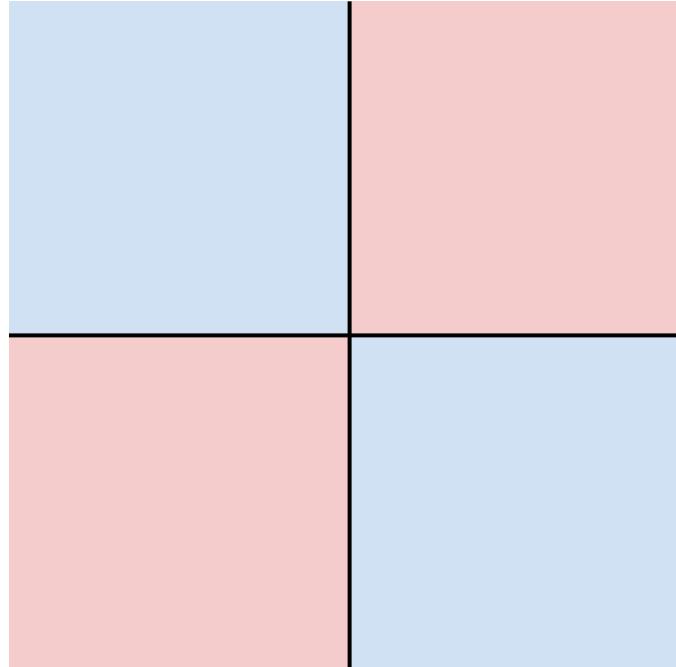
# Hard Cases For A Linear Classifier

**Class 1:**

First and third quadrants

**Class 2:**

Second and fourth quadrants

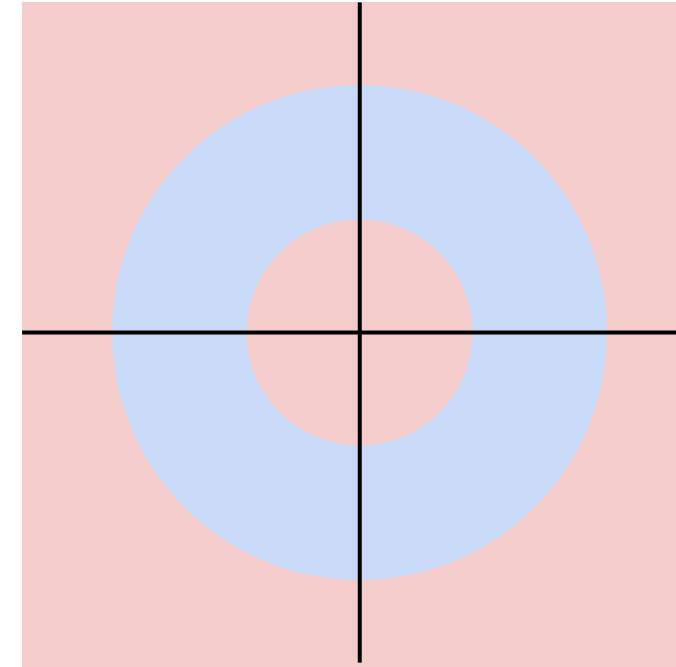


**Class 1:**

$1 \leq L_2 \text{ norm} \leq 2$

**Class 2:**

Everything else

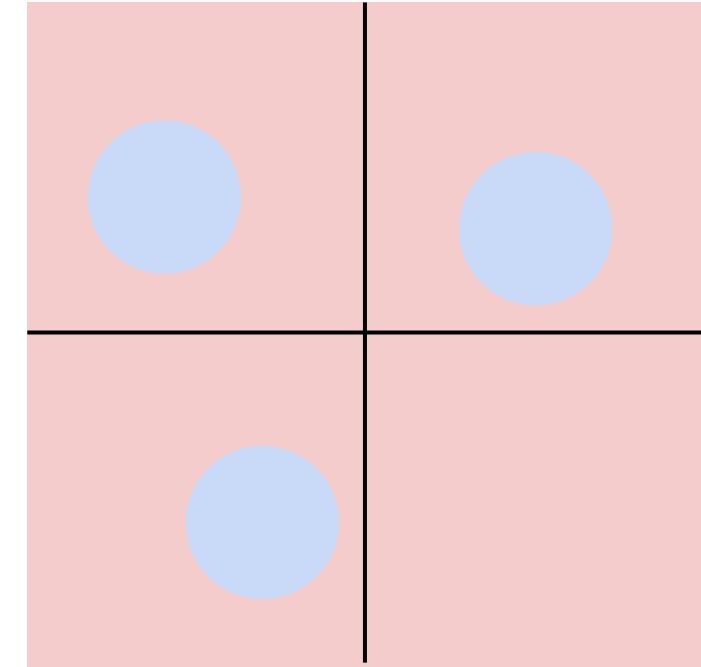


**Class 1:**

Three modes

**Class 2:**

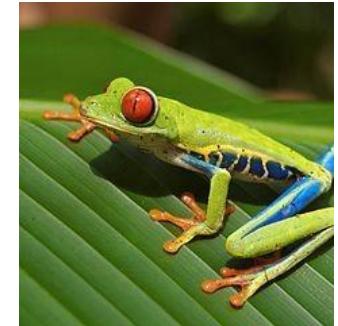
Everything else



# So Far: Defined A (Linear) Score Function $F(x,w) = Wx + B$

Example class scores for 3 images for some  $W$ :

How can we tell whether this  $W$  is good or bad?



airplane	-3.45	-0.51	3.42
automobile	-8.87	<b>6.04</b>	4.64
bird	0.09	5.31	2.65
cat	<b>2.9</b>	-4.22	5.1
deer	4.48	-4.19	2.64
dog	8.02	3.58	5.55
frog	3.78	4.49	<b>-4.34</b>
horse	1.06	-4.37	-1.5
ship	-0.36	-2.09	-4.79
truck	-0.72	-2.93	6.14

[Cat image](#) by [Nikita](#) is licensed under [CC-BY 2.0](#)

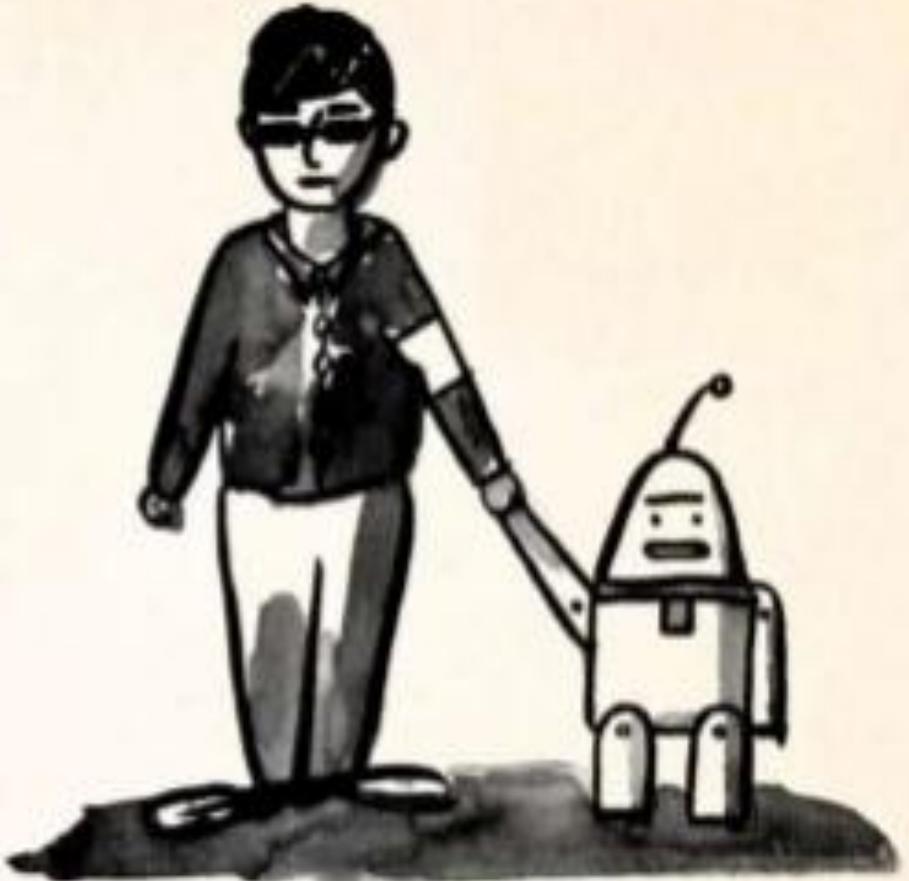
[Car image](#) is [CC0 1.0](#) public domain

[Frog image](#) is in the public domain

$$F(X, W) = WX + B$$

Coming up:

- Loss function  
(quantifying what it means to have a “good” W)
- Optimization  
(start with random W and find a W that minimizes the loss)
- ConvNets!  
(tweak the functional form of f)



**THANK YOU!**  
**QUESTIONS?**